

```
In [3]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
```

```
In [4]: df=pd.read_csv("loan.csv") #read
print(df)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	
0	LP001002	Male	No	0	Graduate	No	\
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	...	...	...	...	...	...	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	
0	5849	0.0	NaN	360.0	\
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

```
In [4]: df.describe() #Display descriptive statistics on the dataset
```

```
Out[4]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
<b>count</b>	614.000000	614.000000	592.000000	600.000000	564.000000
<b>mean</b>	5403.459283	1621.245798	146.412162	342.000000	0.842199
<b>std</b>	6109.041673	2926.248369	85.587325	65.12041	0.364878
<b>min</b>	150.000000	0.000000	9.000000	12.000000	0.000000
<b>25%</b>	2877.500000	0.000000	100.000000	360.000000	1.000000
<b>50%</b>	3812.500000	1188.500000	128.000000	360.000000	1.000000
<b>75%</b>	5795.000000	2297.250000	168.000000	360.000000	1.000000
<b>max</b>	81000.000000	41667.000000	700.000000	480.000000	1.000000

In [5]: `df.head()` *#returns the top five rows by default*

Out[5]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In [7]: `df[df.isna().any(axis=1)]` *#Check if any records in the data have any missing values*

Out[7]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
0	LP001002	Male	No	0	Graduate	No	5849	
11	LP001027	Male	Yes	2	Graduate	NaN	2500	
16	LP001034	Male	No	1	Not Graduate	No	3596	
19	LP001041	Male	Yes	0	Graduate	NaN	2600	
23	LP001050	NaN	Yes	2	Not Graduate	No	3365	
...	...	...	...	...	...	...	...	...
592	LP002933	NaN	No	3+	Graduate	Yes	9357	
597	LP002943	Male	No	NaN	Graduate	No	2987	
600	LP002949	Female	No	3+	Graduate	NaN	416	
601	LP002950	Male	Yes	0	Not Graduate	NaN	2894	
605	LP002960	Male	Yes	0	Not Graduate	No	2400	

134 rows × 13 columns

In [8]: `df.dropna(inplace=True)`

In [9]: `df[df.isna().any(axis=1)]`

Out[9]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
--	---------	--------	---------	------------	-----------	---------------	-----------------	---------

```
In [10]: pip install matplotlib
```

Collecting matplotlibNote: you may need to restart the kernel to use updated packages.

Downloading matplotlib-3.7.1-cp311-cp311-win\_amd64.whl (7.6 MB)

```

0.0/7.6 MB ? eta -:-:--
-- 0.2/7.6 MB 5.3 MB/s eta 0:00:02
--- 0.4/7.6 MB 3.9 MB/s eta 0:00:02
----- 0.7/7.6 MB 4.5 MB/s eta 0:00:02
----- 1.0/7.6 MB 5.0 MB/s eta 0:00:02
----- 1.4/7.6 MB 5.4 MB/s eta 0:00:02
----- 1.5/7.6 MB 5.1 MB/s eta 0:00:02
----- 1.7/7.6 MB 4.9 MB/s eta 0:00:02
----- 1.7/7.6 MB 4.9 MB/s eta 0:00:02
----- 1.7/7.6 MB 4.9 MB/s eta 0:00:02
----- 1.7/7.6 MB 3.8 MB/s eta 0:00:02
----- 2.2/7.6 MB 3.9 MB/s eta 0:00:02
----- 2.4/7.6 MB 4.0 MB/s eta 0:00:02
----- 2.5/7.6 MB 4.1 MB/s eta 0:00:02
----- 2.7/7.6 MB 4.0 MB/s eta 0:00:02
----- 2.9/7.6 MB 4.0 MB/s eta 0:00:02
----- 3.1/7.6 MB 4.0 MB/s eta 0:00:02
----- 3.3/7.6 MB 4.0 MB/s eta 0:00:02
----- 3.5/7.6 MB 4.0 MB/s eta 0:00:02
----- 3.6/7.6 MB 3.9 MB/s eta 0:00:02
----- 3.8/7.6 MB 4.0 MB/s eta 0:00:01
----- 4.0/7.6 MB 3.9 MB/s eta 0:00:01
----- 4.2/7.6 MB 4.0 MB/s eta 0:00:01
----- 4.4/7.6 MB 4.0 MB/s eta 0:00:01
----- 4.5/7.6 MB 4.0 MB/s eta 0:00:01
----- 4.8/7.6 MB 4.0 MB/s eta 0:00:01
----- 5.0/7.6 MB 4.0 MB/s eta 0:00:01
----- 5.1/7.6 MB 3.9 MB/s eta 0:00:01
----- 5.3/7.6 MB 3.9 MB/s eta 0:00:01
----- 5.5/7.6 MB 3.9 MB/s eta 0:00:01
----- 5.7/7.6 MB 3.9 MB/s eta 0:00:01
----- 5.9/7.6 MB 3.9 MB/s eta 0:00:01
----- 6.0/7.6 MB 3.9 MB/s eta 0:00:01
----- 6.2/7.6 MB 3.9 MB/s eta 0:00:01
----- 6.4/7.6 MB 3.9 MB/s eta 0:00:01
----- 6.6/7.6 MB 3.9 MB/s eta 0:00:01
----- 6.8/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.0/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.2/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.4/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.5/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.6/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.6/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.6/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.6/7.6 MB 3.9 MB/s eta 0:00:01
----- 7.6/7.6 MB 3.5 MB/s eta 0:00:00

```

Collecting contourpy>=1.0.1 (from matplotlib)

Downloading contourpy-1.1.0-cp311-cp311-win\_amd64.whl (470 kB)

```

0.0/470.9 kB ? eta -:-:--
----- 297.0/470.9 kB 6.1 MB/s eta 0:00:
01 ----- 460.8/470.9 kB 5.8 MB/s eta 0:00:

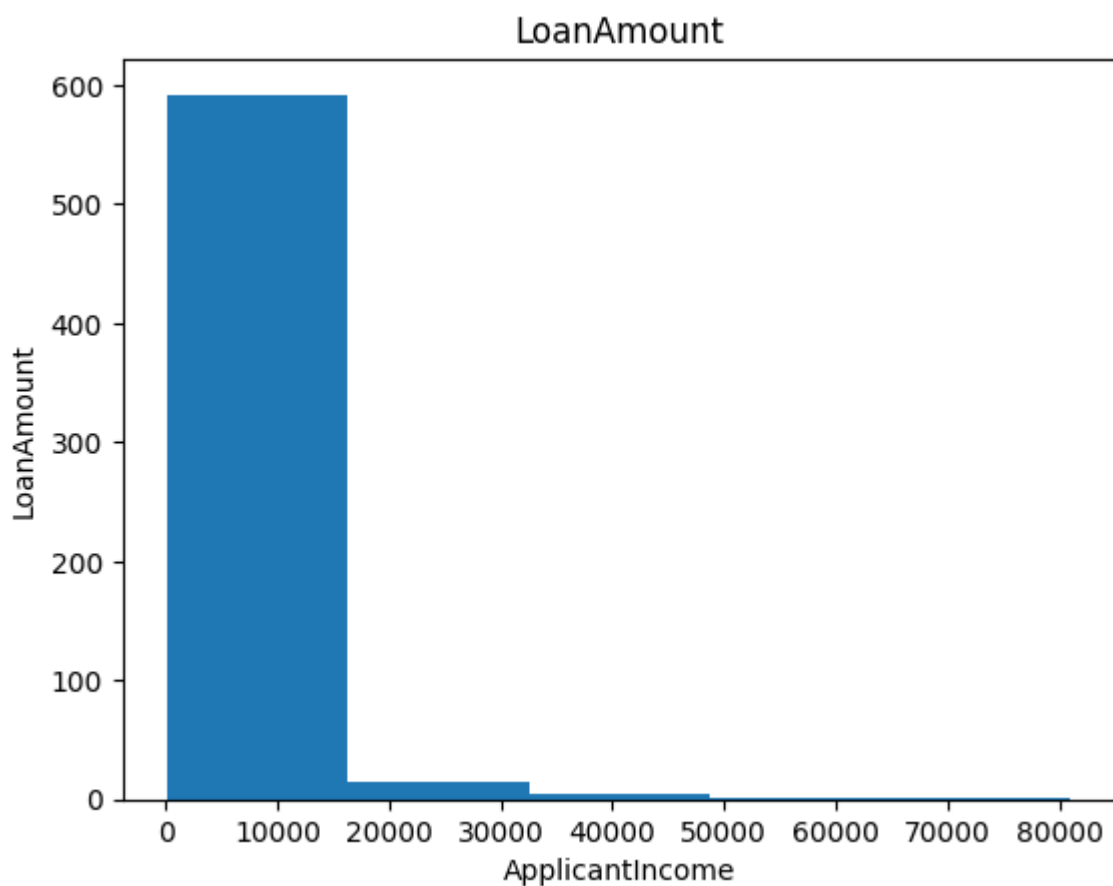
```

```

01 ----- 470.9/470.9 kB 3.7 MB/s eta 0:00:
00
Collecting cycler>=0.10 (from matplotlib)
  Downloading cycler-0.11.0-py3-none-any.whl (6.4 kB)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\user\appdata\loc
al\programs\python\python311\lib\site-packages (from matplotlib) (4.39.4)
Collecting kiwisolver>=1.0.1 (from matplotlib)
  Downloading kiwisolver-1.4.4-cp311-cp311-win_amd64.whl (55 kB)
                                0.0/55.4 kB ? eta -:-:--
----- 51.2/55.4 kB 2.6 MB/s eta 0:00:
01 ----- 55.4/55.4 kB 960.7 kB/s eta 0:00:
00
Requirement already satisfied: numpy>=1.20 in c:\users\user\appdata\local\pro
grams\python\python311\lib\site-packages (from matplotlib) (1.24.3)
Requirement already satisfied: packaging>=20.0 in c:\users\user\appdata\local
\programs\python\python311\lib\site-packages (from matplotlib) (23.0)
Requirement already satisfied: pillow>=6.2.0 in c:\users\user\appdata\local\p
rograms\python\python311\lib\site-packages (from matplotlib) (9.5.0)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Downloading pyparsing-3.0.9-py3-none-any.whl (98 kB)
                                0.0/98.3 kB ? eta -:-:--
----- 92.2/98.3 kB 5.1 MB/s eta 0:00:
01 ----- 98.3/98.3 kB 1.4 MB/s eta 0:00:
00
Requirement already satisfied: python-dateutil>=2.7 in c:\users\user\appdata
\local\programs\python\python311\lib\site-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in c:\users\user\appdata\local\progra
ms\python\python311\lib\site-packages (from python-dateutil>=2.7->matplotlib)
(1.16.0)
Installing collected packages: pyparsing, kiwisolver, cycler, contourpy, matp
lotlib
Successfully installed contourpy-1.1.0 cycler-0.11.0 kiwisolver-1.4.4 matplot
lib-3.7.1 pyparsing-3.0.9


```

```
In [13]: import matplotlib.pyplot as plt #Build a graph visualizing the distribution of
import pandas as pd
df=pd.read_csv("loan.csv")
plt.hist(df['ApplicantIncome'], bins=5)
plt.xlabel('ApplicantIncome')
plt.ylabel('LoanAmount')
plt.title('LoanAmount')
plt.show()
```



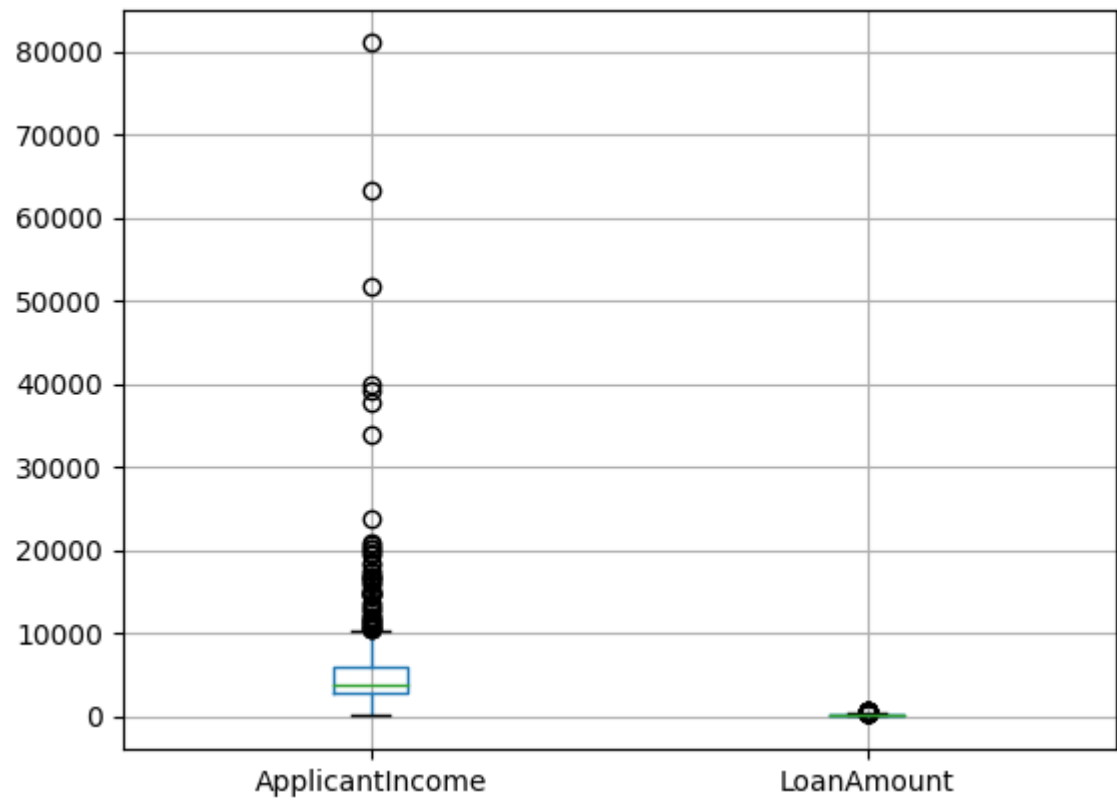
```
In [15]: x=df["ApplicantIncome"].to_numpy()
y=df["LoanAmount"].to_numpy()
df1=plt.scatter(x,y)
plt.title("relation visualization with scatter plot")
plt.xlabel("ApplicantIncome")
plt.ylabel("LoanAmount")
```

Out[15]: Text(0, 0.5, 'LoanAmount')



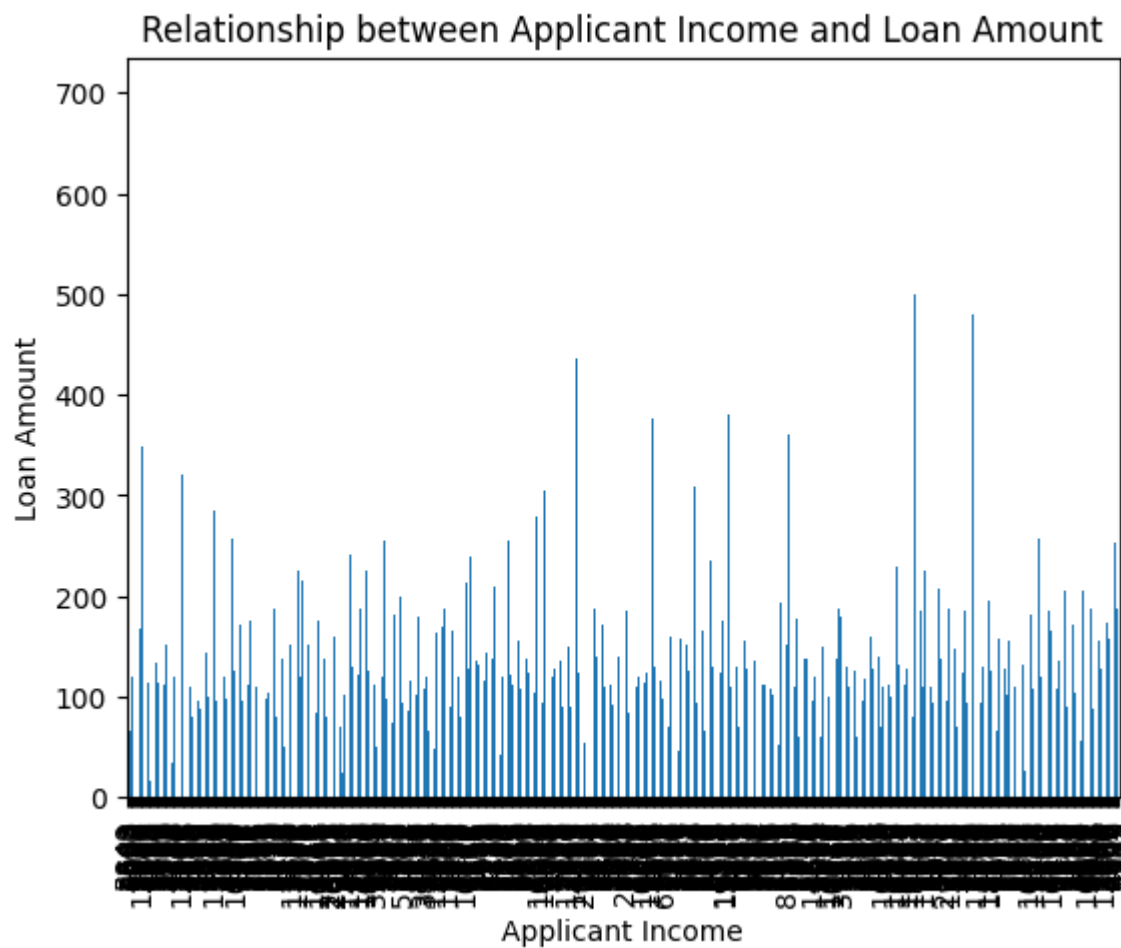
```
In [19]: df.boxplot(column=['ApplicantIncome', 'LoanAmount']) #Build a graph visualizing
```

```
Out[19]: <Axes: >
```





```
In [6]: df.plot(kind='bar', x='ApplicantIncome', y='LoanAmount', legend=False) #without  
plt.xlabel('Applicant Income')  
plt.ylabel('Loan Amount')  
plt.title('Relationship between Applicant Income and Loan Amount')  
plt.show()
```



```
In [5]: import seaborn as sns
```

In [16]: `pip install seaborn`

Collecting seabornNote: you may need to restart the kernel to use updated packages.

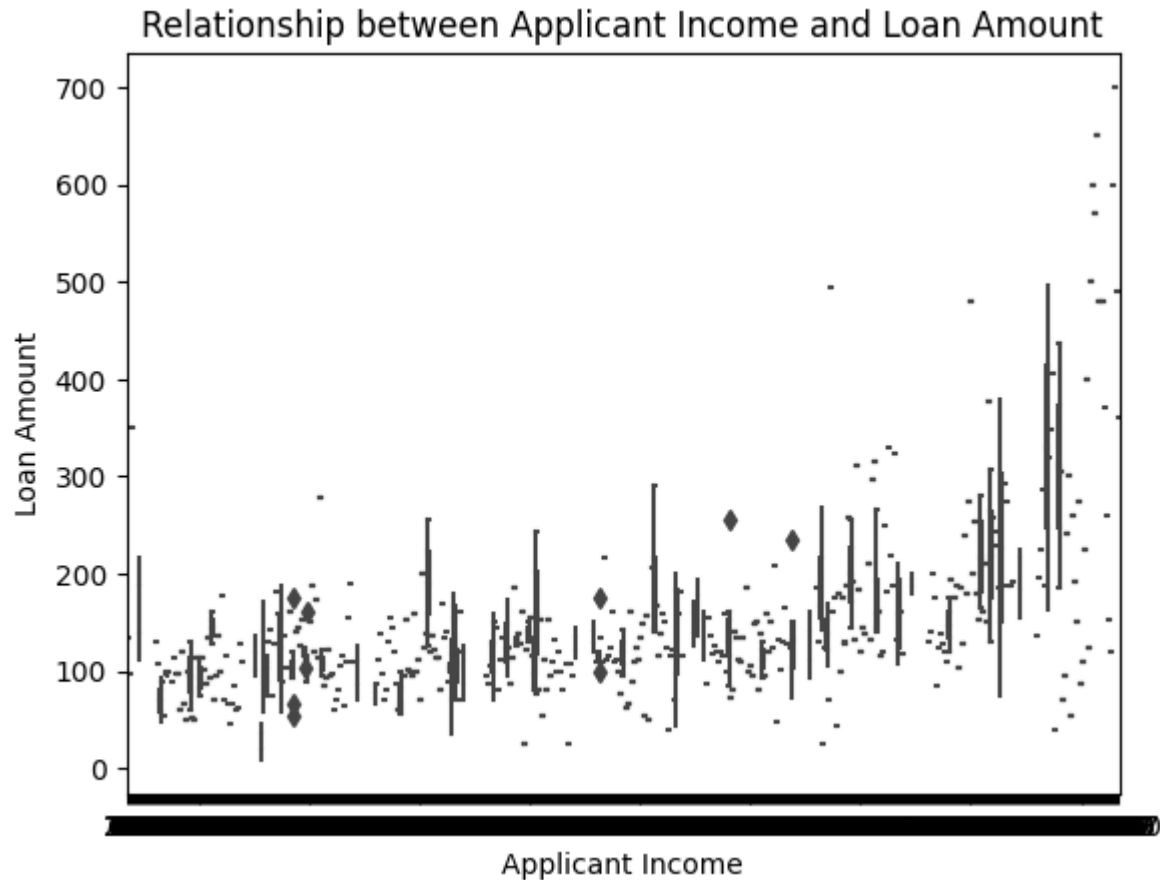
```

    Downloading seaborn-0.12.2-py3-none-any.whl (293 kB)
      0.0/293.3 kB ? eta -:--:--
-----
01      204.8/293.3 kB 6.3 MB/s eta 0:00:
01      286.7/293.3 kB 3.5 MB/s eta 0:00:
01      293.3/293.3 kB 2.0 MB/s eta 0:00:
00
Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (1.24.3)
Requirement already satisfied: pandas>=0.25 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (2.0.1)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from seaborn) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.1.0)
Requirement already satisfied: cycler>=0.10 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.39.4)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.0)
Requirement already satisfied: pillow>=6.2.0 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.5.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)
Requirement already satisfied: tzdata>=2022.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)
Requirement already satisfied: six>=1.5 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)
Installing collected packages: seaborn
Successfully installed seaborn-0.12.2

```

```
In [17]: import seaborn as sns
```

```
In [18]: sns.boxplot(x=df['ApplicantIncome'], y=df['LoanAmount']) #using seaborn
plt.xlabel('Applicant Income')
plt.ylabel('Loan Amount')
plt.title('Relationship between Applicant Income and Loan Amount')
plt.show()
```



```
In [19]: unique_values = df['Property_Area'].unique() #Display unique values of a categorical variable
print(unique_values)
```

```
['Urban' 'Rural' 'Semiurban']
```

```
In [20]: unique_values = df['Education'].unique() #Display unique values of a categorical variable
print(unique_values)
```

```
['Graduate' 'Not Graduate']
```

In [24]: `pip install scipy`

```
Collecting scipy
  Downloading scipy-1.10.1-cp311-cp311-win_amd64.whl (42.2 MB)
    0.0/42.2 MB ? eta -:--:--
    0.1/42.2 MB 3.2 MB/s eta 0:0
0:14    0.2/42.2 MB 2.9 MB/s eta 0:0
0:15    0.3/42.2 MB 2.6 MB/s eta 0:0
0:17    0.3/42.2 MB 1.8 MB/s eta 0:0
0:23    0.5/42.2 MB 2.2 MB/s eta 0:0
0:19    0.6/42.2 MB 2.3 MB/s eta 0:0
0:19    0.7/42.2 MB 2.1 MB/s eta 0:0
0:20    0.7/42.2 MB 2.0 MB/s eta 0:0
0:21    0.8/42.2 MB 2.0 MB/s eta 0:0
```

In [22]: `import scipy.stats as stats`

In [24]: `# Build a contingency table of two potentially related categorical variables. (`  
`contingency_table = pd.crosstab(df['Property_Area'], df['Self_Employed'])`  
`print(contingency_table)`  
  
`# Perform the chi-square test of independence`  
`chi2, p_value, _, _ = stats.chi2_contingency(contingency_table)`  
`print('Chi-square statistic:', chi2)`  
`print('p-value:', p_value)`

```
Self_Employed    No  Yes
Property_Area
Rural            143   26
Semiurban        191   32
Urban            166   24
Chi-square statistic: 0.5803134707599138
p-value: 0.7481462973944264
```

```
In [25]: # Retrieve one or more subset of rows based on two or more criteria and present
subset1 = df[(df['Gender'] == 'Male') & (df['Married'] == 'Yes')]
subset2 = df[(df['Loan_Status'] == 'N') & (df['Property_Area'] == 'Rural')]

# Present descriptive statistics on the subset(s)
subset1_stats = subset1.describe()
subset2_stats = subset2.describe()

print("Subset 1 descriptive statistics:")
print(subset1_stats)

print("\nSubset 2 descriptive statistics:")
print(subset2_stats)
```

Subset 1 descriptive statistics:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
count	357.000000	357.000000	343.000000	348.000000 \
mean	5529.540616	1828.330308	154.011662	335.931034
std	6743.209021	2096.367198	83.025254	67.342095
min	150.000000	0.000000	17.000000	12.000000
25%	2882.000000	0.000000	108.000000	360.000000
50%	3875.000000	1619.000000	132.000000	360.000000
75%	5829.000000	2500.000000	180.000000	360.000000
max	81000.000000	20000.000000	600.000000	480.000000

	Credit_History
count	326.000000
mean	0.846626
std	0.360902
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

Subset 2 descriptive statistics:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
count	69.000000	69.000000	66.000000	67.000000 \
mean	6497.521739	1436.507246	158.742424	345.134328
std	9949.733021	1618.225813	96.412305	56.751691
min	150.000000	0.000000	46.000000	84.000000
25%	3365.000000	0.000000	110.500000	360.000000
50%	4283.000000	1287.000000	133.000000	360.000000
75%	6216.000000	2200.000000	172.750000	360.000000
max	81000.000000	5302.000000	570.000000	480.000000

	Credit_History
count	65.000000
mean	0.600000
std	0.49371
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

In [1]: `pip install scikit-learn`

Requirement already satisfied: scikit-learn in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (1.2.2)  
 Requirement already satisfied: numpy>=1.17.3 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from scikit-learn) (1.24.3)  
 Requirement already satisfied: scipy>=1.3.2 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from scikit-learn) (1.10.1)  
 Requirement already satisfied: joblib>=1.1.1 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from scikit-learn) (1.3.0)  
 Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\user\appdata\local\programs\python\python311\lib\site-packages (from scikit-learn) (3.1.0)

Note: you may need to restart the kernel to use updated packages.

In [1]: `import numpy as np` *#Conduct a statistical test of the significance of the*  
`import scipy.stats as stats`

*# Define the two subsets of data*

`data_a = np.array([5849, 4583, 3000, 2583, 6000, 5417, 2333, 3036, 4006])`

`data_b = np.array([0, 128, 66, 120, 141, 267, 95, 158, 168])`

*# Conduct the independent t-test*

`t_statistic, p_value = stats.ttest_ind(data_a, data_b)`

*# Print the results*

`print("t-statistic:", t_statistic)`

`print("p-value:", p_value)`

t-statistic: 8.289381635365402

p-value: 3.486682839706522e-07

In [9]: *#Create one or more tables that group the data by a certain categorical variable*

*# Group the data by 'Gender' and calculate the mean and sum for selected columns*  
`grouped_gender = df.groupby('Gender').agg({'ApplicantIncome': 'mean', 'CoapplicantIncome': 'sum'})`

*# Group the data by 'Education' and calculate the mean and sum for selected columns*  
`grouped_education = df.groupby('Education').agg({'ApplicantIncome': 'mean', 'CoapplicantIncome': 'sum'})`

*# Print the grouped data for 'Gender'*

`print("Grouped Data by Gender:")`

`print(grouped_gender)`

Grouped Data by Gender:

	ApplicantIncome	CoapplicantIncome	LoanAmount
Gender			
Female	4593.954128	1138.504587	13810.0
Male	5455.895745	1764.442383	70155.0

```

In [7]: from sklearn.model_selection import train_test_split    #Implement a linear regression
from sklearn.linear_model import LinearRegression
# Remove rows with missing values
df.dropna(subset=['LoanAmount'], inplace=True)

# Split the data into training and testing sets
X = df[['ApplicantIncome']]
y = df['LoanAmount']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Create and fit the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Plot the actual vs. predicted values
plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, y_pred, color='red', linewidth=2, label='Predicted')
plt.xlabel('Applicant Income')
plt.ylabel('Loan Amount')
plt.title('Linear Regression')
plt.legend()
plt.show()

# Print the model coefficients and intercept
print('Coefficients:', model.coef_)
print('Intercept:', model.intercept_)

```



Coefficients: [0.00823739]  
Intercept: 103.70294803264798

In [ ]: